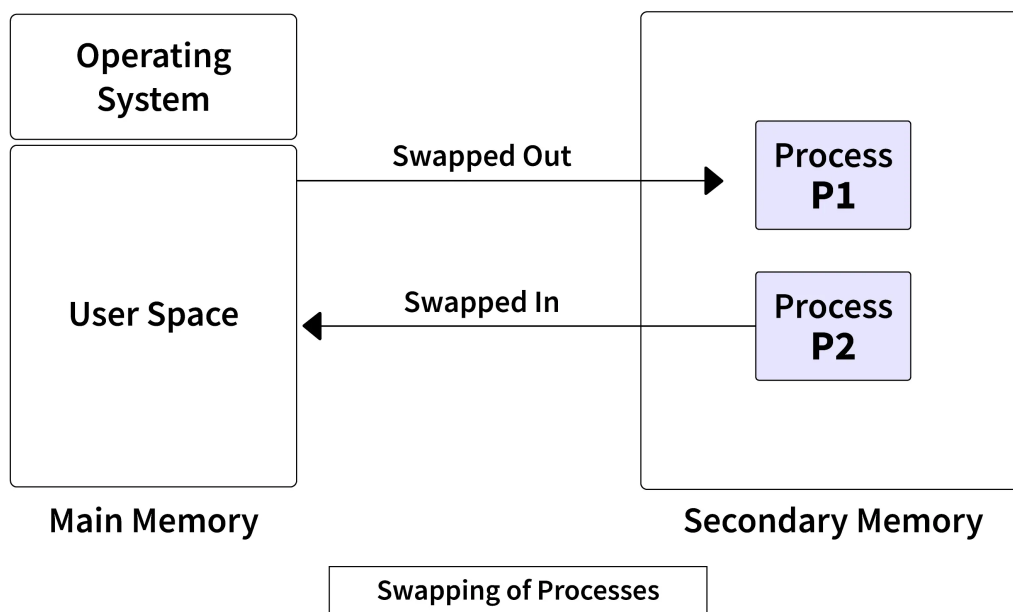


1. Write a short note on swapping with neat diagram ?

Swapping in Operating Systems

When a process is executed it must have resided in memory. Swapping is a process of swapping a process temporarily into secondary memory from the main memory, which is fast compared to secondary memory. Swapping allows more processes to be run and can be fit into memory at one time. The main part of swapping is transferred time and the total time is directly proportional to the amount of memory swapped. Swapping is also known as roll-out, or roll because if a higher priority process arrives and wants service, the memory manager can swap out the lower priority process and then load and execute the higher priority process. After finishing higher priority work, the lower priority process swapped back in memory and continued to the execution process.



SCALER
Topics

Advantages

- If there is low main memory so some processes may have to wait for much long but by using swapping process do not have to wait long for execution on CPU.
- It utilizes the main memory.

- Using only single main memory, multiple processes can be run by CPU using swap partition.
- The concept of virtual memory starts from here and it utilizes it in a better way.
- This concept can be useful in priority-based scheduling to optimize the swapping process.

Disadvantages

- If there is low main memory resource and the user is executing too many processes and suddenly the power of the system goes off there might be a scenario where data gets erased of the processes which are took parts in swapping.
- Chances of a number of page faults occur
- Low processing performance

Example

Suppose the user process's size is 2048KB and is a standard hard disk where swapping has a data transfer rate of 1Mbps. Calculate how long it will take to transfer from main memory to secondary memory.

- User process size is 2048Kb
- Data transfer rate is 1Mbps = 1024 kbps
- Time = process size / transfer rate
- Time = $2048 / 1024 = 2$ seconds or 2000 milliseconds

Now taking swap-in and swap-out time, the process will take 4000 ms or 4 seconds.

2. What is paging? Explain its structure for 32-byte memory with 4-byte pages?

Paging in Memory Management

Paging is a memory management scheme that eliminates the need for a contiguous allocation of physical memory. This scheme permits the physical address space of a process to be non-contiguous. The mapping from virtual to physical address is done by the memory management unit (MMU) which is a hardware device and this mapping is known as the paging technique.

- The Physical Address Space is conceptually divided into several fixed-size blocks, called frames.
- The Logical Address Space is also split into fixed-size blocks, called pages.

- Page Size = Frame Size

The address generated by the CPU is divided into:

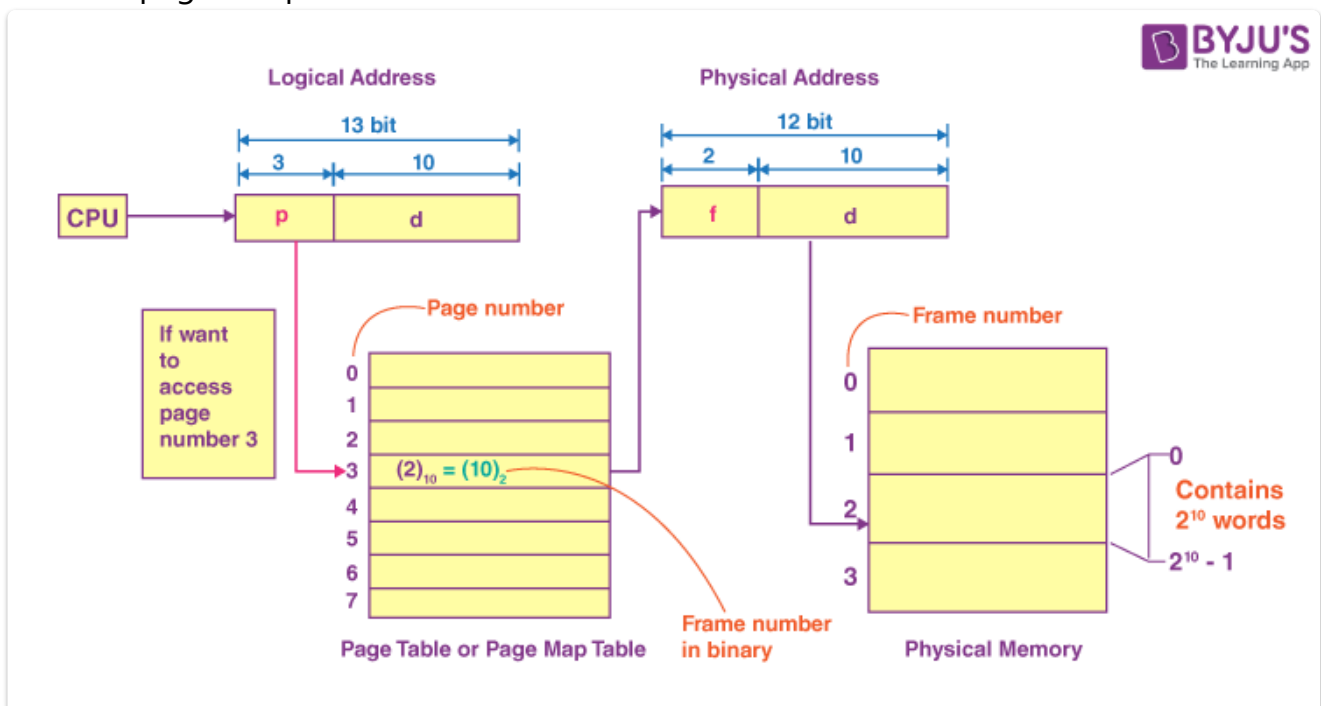
- Page Number (p): Number of bits required to represent the pages in Logical Address Space or Page number
- Page Offset (d): Number of bits required to represent a particular word in a page or page size of Logical Address Space or word number of a page or page offset.

Physical Address is divided into:

- Frame Number (f): Number of bits required to represent the frame of Physical Address Space or Frame number frame
- Frame Offset (d): Number of bits required to represent a particular word in a frame or frame size of Physical Address Space or word number of a frame or frame offset.

Structure of Paging:

In a paging system, the physical address space is divided into fixed-size blocks called frames, and the logical address space is divided into fixed-size blocks called pages. The size of a page is equal to the size of a frame.



For example, let's consider a memory system with a total of 32 bytes and 4-byte pages:

- **Memory Size:** 32 bytes
- **Page Size:** 4 bytes
- **Number of Frames:** Total memory size / Page size = 32 bytes / 4 bytes per page = 8 frames

Page Table:

Each process has its own page table, which maps logical page numbers to physical frame numbers. Initially, the frames are empty, and the pages of the processes are stored contiguously in memory.

Process Status:

Processes P1, P2, P3, and P4 are initially loaded into memory, with their pages mapped to frames as follows:

Process	Page Number	Frame Number
P1	0	0
P1	1	1
P1	2	2
P1	3	3
P2	0	4
P2	1	5
P2	2	6
P2	3	7
P3	0	0
P3	1	1
P3	2	2
P3	3	3
P4	0	4
P4	1	5
P4	2	6
P4	3	7

Process State Change:

Processes P2 and P4 are shifted to the waiting state, leaving 8 empty frames in memory.

3. What is File system? Explain different file system attributes and file system operations?

A file system is a method used by operating systems to organize and manage files on a storage device, such as a hard drive or solid-state drive. It provides a structure for

storing, accessing, and managing files efficiently. Below are explanations of different file system attributes and file system operations:

Attributes of a File:

1. **Name:** Every file has a unique name by which it is recognized in the file system. Names must be unique within a directory.
2. **Identifier/Extension:** Files may have extensions that identify their type, such as .txt for text files, .mp4 for video files, etc.
3. **Type:** Files are classified into different types based on their format or content, such as text files, audio files, executable files, etc.
4. **Location:** Files are stored at specific locations within the file system, which is managed by the operating system.
5. **Size:** The size of a file indicates the amount of storage space it occupies in the file system, usually measured in bytes.
6. **Protection:** Files may have different levels of access permissions assigned to them, such as read-only, write-only, or execute permissions.
7. **Time and Date:** Each file carries a timestamp indicating the time and date of its creation, modification, or access.

File System Operations:

1. **Create:** This operation is used to create a new file in the file system. The file system allocates space for the new file and adds an entry to the directory structure.
2. **Open:** Opening a file allows for reading from or writing to it. The operating system invokes the open system call and passes the file name to the file system.
3. **Write:** Writing to a file involves adding or modifying data within the file. The system call specifies the file name and the data to be written.
4. **Read:** Reading from a file retrieves data from the file. A read pointer maintained by the operating system tracks the position up to which data has been read.
5. **Re-position or Seek:** This operation moves the file pointer to a specific position within the file, allowing for random access.
6. **Delete:** Deleting a file removes it from the file system, freeing up disk space occupied by the file. The file system searches for the file's directory entry and releases associated resources.
7. **Truncate:** Truncating a file removes its contents while retaining its attributes. The file is not completely deleted, but the data inside it is replaced.
8. **Close:** Closing a file finalizes changes made during file operations and releases resources associated with the file.

9. **Append:** Appending data to a file adds new information to the end of the file without altering existing content.
10. **Rename:** Renaming a file changes its name while preserving its contents and attributes. The directory entry for the file is updated with the new name.

4. What is segmentation? Explain its structure for 32-byte memory with 4-byte segmentation?

Segmentation is a memory management technique where a process is divided into segments, which are logical units such as main programs, procedures, functions, methods, objects, local variables, global variables, etc. Segmentation provides a user's view of the process, mapping logical addresses to physical memory.

Here is the structure of segmentation for a 32-byte memory with 4-byte segmentation:

1. Segment Table:


- The segment table is a data structure that stores information about each segment of the process.
- It maps a two-dimensional logical address (segment number, offset) to a one-dimensional physical address.
- Each entry in the segment table contains:
 - Base Address: The starting physical address where the segment resides in memory.
 - Segment Limit: Specifies the length of the segment.

2. Address Generation:

- The address generated by the CPU is divided into two parts:
 - Segment Number (s): Number of bits required to represent the segment.
 - Segment Offset (d): Number of bits required to represent the size of the segment.

3. The Segment number is mapped to the segment table. The limit of the respective segment is compared with the offset. If the offset is less than the limit then the address is valid otherwise it throws an error as the address is invalid. In the case of valid addresses, the base address of the segment is added to the offset to get the physical address of the actual word in the main memory.

Logical View of Segmentation




Logical Address Space

Segment Number

	base address	Limit
0	500	600
1	2500	800
2	1500	400
3	4600	200
4	3800	400

Segment Table



Physical Address Space

Example for 32-Byte Memory with 4-Byte Segmentation:

1. **Memory Size:** The total memory capacity is 32 bytes.
2. **Segment Size:** Each segment is 4 bytes long, defining the granularity of memory allocation and addressing.
3. **Segment Table:**
 - Contains 8 entries (for 8 segments in total).
 - Each entry includes:
 - Segment Number (s): Identifies the segment.
 - Base Address: Starting physical address of the segment.
 - Segment Limit: Specifies the length of the segment (4 bytes each).

Segment Number	Base Address	Limit
0	0	4
1	4	4
2	8	4
3	12	4

Segment Number	Base Address	Limit
4	16	4
5	20	4
6	24	4
7	28	4

4. Address Generation:

- CPU generates logical addresses consisting of:
 - Segment Number (s): Indicates the segment to access.
 - Segment Offset (d): Specifies the byte's location within the segment (0 to 3).

5. Address Translation and Validation:

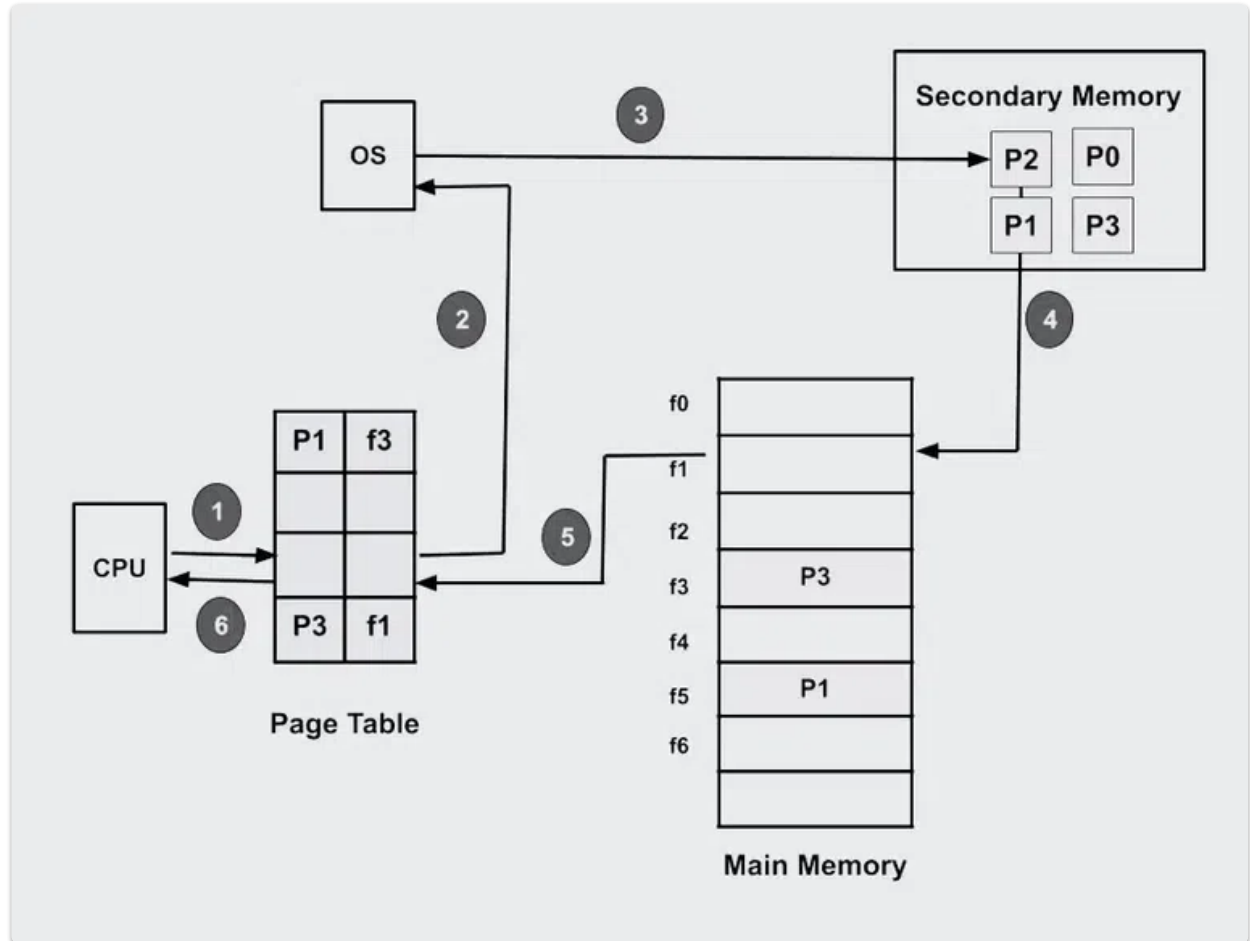
- Logical address translated into physical address by adding segment base address to segment offset.
- Before translation, offset is validated against segment limit to ensure it doesn't exceed the segment boundary, preventing access violations.

5. Write a short note on Demand paging with neat diagram?

Demand paging is a memory management technique used by operating systems to efficiently utilize memory resources. It involves bringing only necessary pages of a process into memory when they are needed, rather than loading the entire process at once. Demand paging helps reduce the amount of physical memory needed and decreases the time required for swapping processes in and out of memory.

1. **Page Faults:** When a process attempts to access a page that is not currently in memory, a page fault occurs. The operating system then brings the required page into memory from disk.
2. **Valid-Invalid Bit:** Hardware support is needed to distinguish between pages that are in memory and those that are on disk. This is achieved using a valid-invalid bit scheme. A page marked as valid is in memory, while a page marked as invalid is on disk.
3. **Handling Page Faults:** When a page fault occurs, the operating system follows specific steps:
 - Verify the memory access request made by the process.
 - If the request is valid, locate a free frame in memory.
 - Schedule an I/O operation to move the required page from disk to memory.

- Update the process's page table with the new frame number and mark the page as valid.
 - Restart the instruction that caused the page fault.
4. **Dirty Bit:** A dirty bit is a flag that indicates whether a page in memory has been modified since it was brought into memory. This helps the operating system determine which pages need to be written back to disk when they are replaced.



Advantages of Demand Paging:

5. Allows for large virtual memory sizes.
6. More efficient use of physical memory.
7. Supports unconstrained multiprogramming, allowing a higher degree of multiprogramming without strict limits on address space size.

Disadvantages of Demand Paging:

1. Increased overhead due to managing page tables and handling page faults.
2. Lack of explicit constraints on job address space size can lead to inefficient memory usage.

6. Explain different Page Replacement Algorithms with example?

Page replacement algorithms are crucial components of memory management in operating systems, especially in systems that use paging. Here are explanations of different page replacement algorithms along with examples:

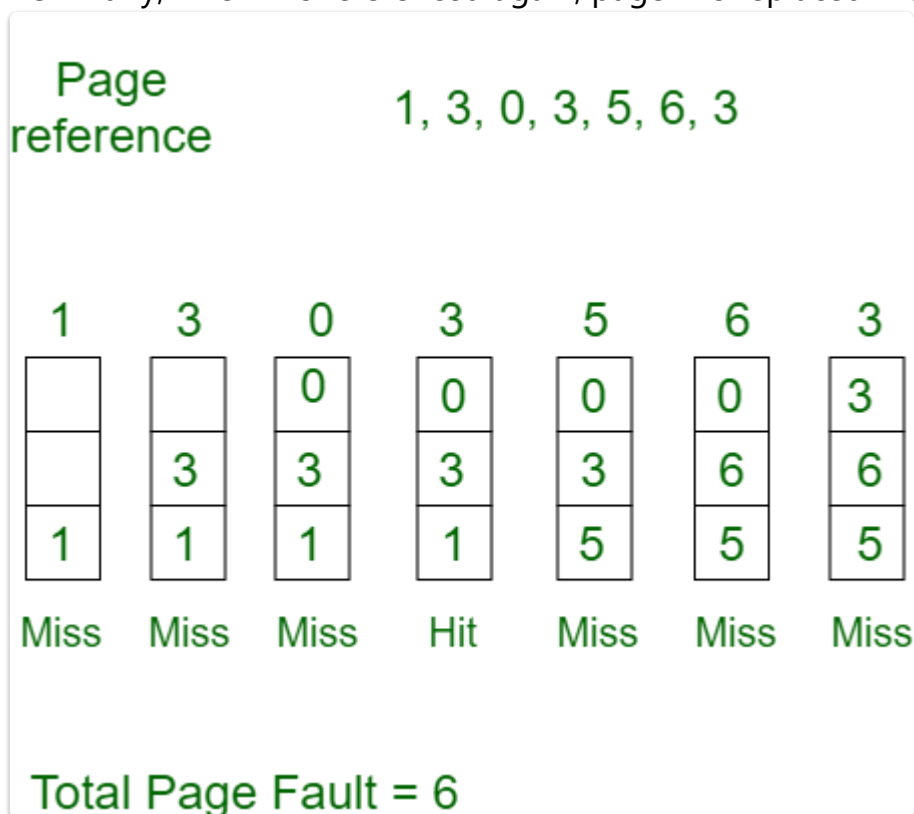
1. First-In-First-Out (FIFO) Algorithm:

- In the FIFO algorithm, the oldest page in main memory is selected for replacement.
- It is easy to implement and involves maintaining a list of pages in the order they were brought into memory.
- Pages are replaced from the tail of the list, and new pages are added to the head.

- Example:

Consider a memory with a capacity of 3 pages (A, B, C) and the reference string: 1, 2, 3, 4, 1, 2

- Initially, A, B, and C are loaded into memory.
- When 4 is referenced, page A (the oldest) is replaced with 4.
- Similarly, when 1 is referenced again, page B is replaced with 1.



2. Optimal Page Replacement Algorithm:

- The optimal page replacement algorithm replaces the page that will not be used for the longest period of time in the future.
- It requires knowledge of future memory references, which is not practical in real systems.
- The optimal algorithm serves as a benchmark for comparing other algorithms.

- Consider a memory with a capacity of 3 pages and the reference string: 1, 2, 3, 4, 1, 2
- When 4 is referenced, page A or B or C can be replaced, but it's not possible to predict which one will not be used in the future.

Page reference: 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 3

No. of Page frame - 4

7	0	1	2	0	3	0	4	2	3	0	3	2	3
			2	2	2	2	2	2	2	2	2	2	2
		1	1	1	1	1	4	4	4	4	4	4	4
	0	0	0	0	0	0	0	0	0	0	0	0	0
7	7	7	7	7	3	3	3	3	3	3	3	3	3
Miss	Miss	Miss	Miss	Hit	Miss	Hit	Miss	Hit	Hit	Hit	Hit	Hit	Hit

Total Page Fault = 6

3. Least Recently Used (LRU) Algorithm:

- The LRU algorithm replaces the page that has not been used for the longest time in main memory.
- It requires maintaining a list of pages and updating their access timestamps.
- The page with the oldest timestamp is selected for replacement.
- Example:

Consider a memory with a capacity of 3 pages and the reference string: 1, 2, 3, 4, 1, 2

- Initially, A, B, and C are loaded into memory.
- When 4 is referenced, the least recently used page, A, is replaced with 4.
- When 1 is referenced again, the least recently used page, B, is replaced with 1.

Page reference 7,0,1,2,0,3,0,4,2,3,0,3,2,3 **No. of Page frame - 4**

7	0	1	2	0	3	0	4	2	3	0	3	2	3
			2	2	2	2	2	2	2	2	2	2	2
		1	1	1	1	1	4	4	4	4	4	4	4
	0	0	0	0	0	0	0	0	0	0	0	0	0
7	7	7	7	7	3	3	3	3	3	3	3	3	3
Miss	Miss	Miss	Miss	Hit	Miss	Hit	Miss	Hit	Hit	Hit	Hit	Hit	Hit

Total Page Fault = 6

Here LRU has same number of page fault as optimal but it may differ according to question.

7. List out various methods for accessing the file and explain them?

Accessing files in a computer system involves a variety of methods and system calls aimed at creating, reading, writing, and manipulating file data and attributes. Here's a comprehensive explanation of the methods outlined:

1. **create()**: The `create()` system call is responsible for creating a new file in the file system. This operation involves finding available space in the file system and adding an entry for the new file in the directory.
2. **open()**: When a file needs to be accessed, the `open()` system call is used to open the file for reading, writing, or both. This operation adds the file's entry to the open file table and maintains an open count associated with each file, enabling tracking of the number of processes with the file open.
3. **read()**: The `read()` system call allows data to be read from a file. It involves specifying the file name and read pointer to determine the read location within the file. After reading, the read pointer is updated to reflect the next position to read from.
4. **write()**: Conversely, the `write()` system call is used to write data to a file. Similar to `read()`, it requires specifying the file name and the data to be written. The file system searches the directory for the file's location and updates the write pointer to indicate where the next write operation should occur.
5. **close()**: Once file operations are complete, the `close()` system call is used to close the file. This operation decrements the open count associated with the file and closes the file when the open count reaches zero, releasing associated resources.
6. **lseek()**: The `lseek()` system call is employed to reposition the file offset of an open file descriptor. It facilitates random access within a file by moving the file pointer to a specified position, thus enabling efficient navigation and manipulation of file data.
7. **stat()**: The `stat()` system call retrieves file attributes, such as size, permissions, timestamps, and other details, for a given file path. This information is crucial for understanding and managing file properties within the system.
8. **ioctl()**: For special files and devices, the `ioctl()` system call provides input/output control. It enables various operations on devices, including setting parameters, querying device status, and performing specialized I/O operations beyond standard file access.
9. **delete()**: The `delete()` operation is used to remove a file from the file system entirely. It involves searching the directory for the named file, releasing its

associated space, and erasing its directory entry, effectively eliminating the file from the system.

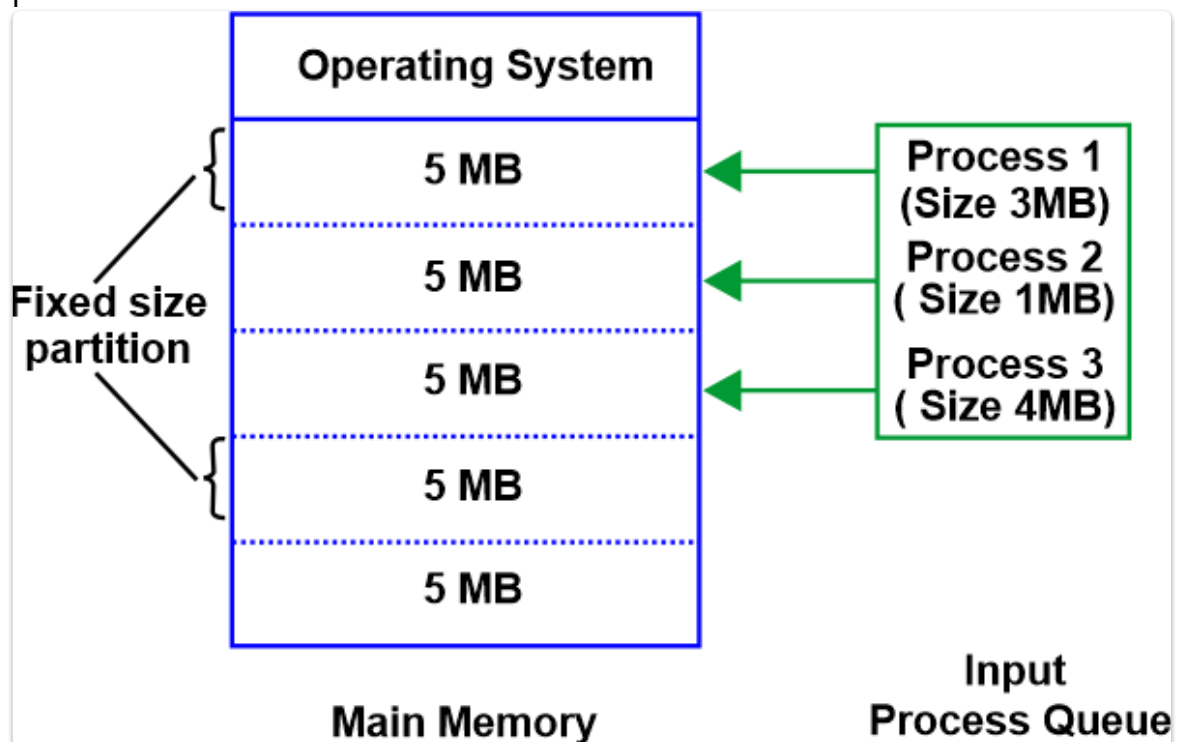
10. **seek() (Reposition)**: Similar to `lseek()`, the `seek()` operation repositions the current file position pointer within a file, allowing for efficient navigation and manipulation of file data.

8. Explain different memory allocation techniques with diagram?

Memory Allocation Techniques

1. Fixed Partition Allocation


- **Description**: Memory is divided into fixed-sized partitions, each capable of holding one process.
- **Operation**:
 - Memory is initially divided into fixed-sized partitions.
 - Each partition holds one process.
 - When a process arrives, it is allocated to the smallest partition that can accommodate it.
 - After the process completes, the partition becomes available for the next process.



2. Dynamic Partition Allocation (Variable Partitioning)

- **Description**: Memory is divided into variable-sized partitions, each capable of holding one process.

- **Operation:**
 - Memory starts as one large block.
 - When a process arrives, the OS searches for a suitable-sized partition.
 - The process is allocated memory in the selected partition.
 - After the process completes, the partition becomes available for future allocations.



3. First Fit Allocation

- **Description:** Memory is allocated to the first available partition large enough to hold the process.
- **Operation:**
 - The OS allocates memory to the first partition that can accommodate the process.

4. Best Fit Allocation

- **Description:** Memory is allocated to the smallest available partition large enough to hold the process.
- **Operation:**
 - The OS allocates memory to the smallest partition that minimizes wasted space.

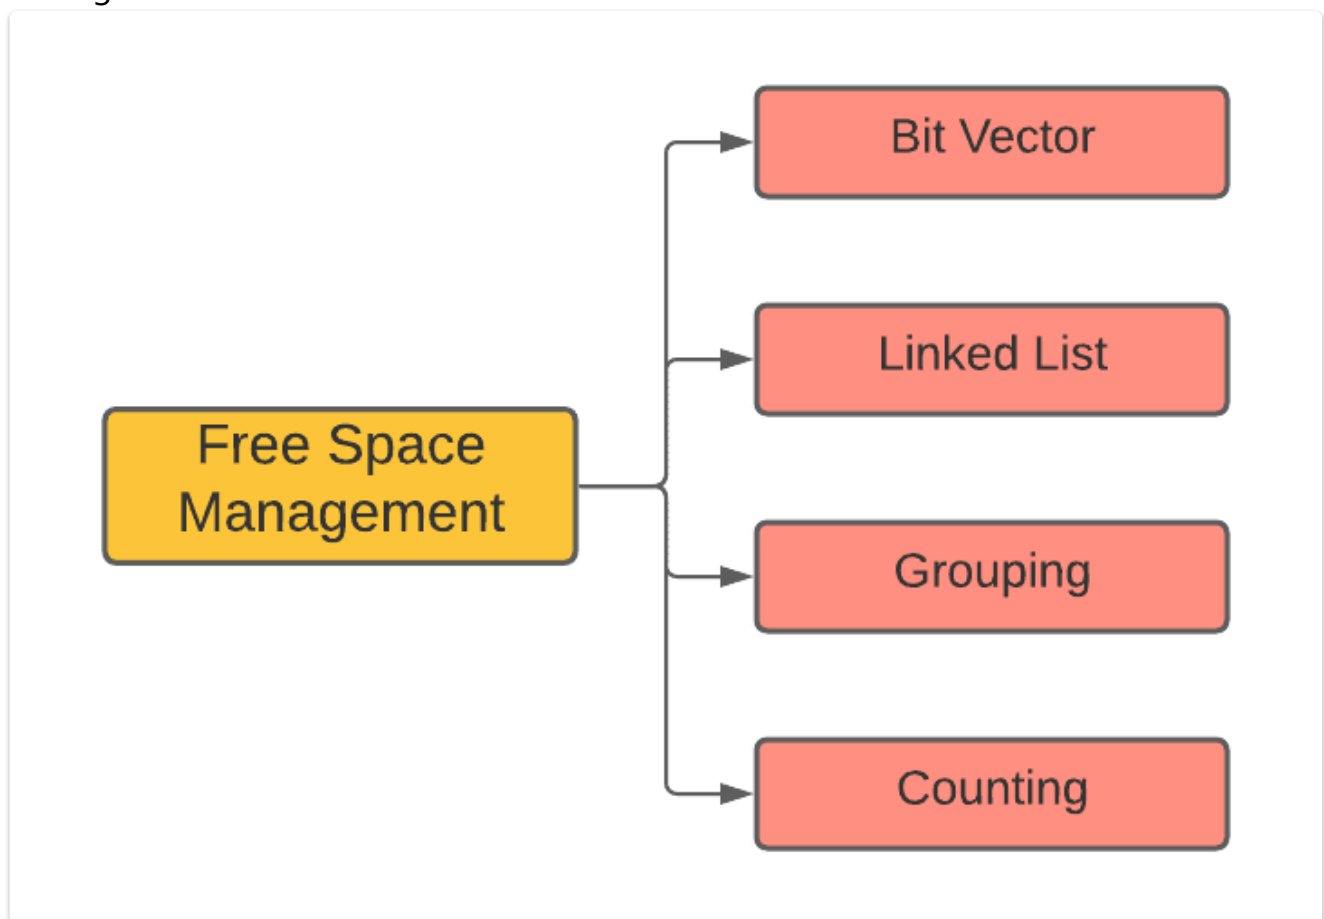
5. Worst Fit Allocation

- **Description:** Memory is allocated to the largest available partition large enough to hold the process.
- **Operation:**
 - The OS allocates memory to the largest partition, potentially leaving larger holes for future allocations.

9. List out various methods for Free-space management and explain them?

Free-Space Management Techniques


Free-space management refers to the techniques used by a file system to keep track of available space on the disk for storing files. Here are various methods for free-space management:



1. Bitmap or Bit Vector:

- A bitmap or bit vector is a collection of bits where each bit corresponds to a disk block.
- The bit value indicates whether the block is allocated (0) or free (1).
- It is simple to understand and efficient for finding the first free block by scanning the bitmap.

- The drawback is that it may require a significant amount of memory to store the bitmap, especially for large disk sizes.



SCALER
Topics

2. Linked List (Free List):

- In this method, free disk blocks are linked together using pointers.
- Each free block contains a pointer to the next free block, forming a linked list.
- The advantage is that there is no waste of space, and finding the first free block is efficient.
- However, getting contiguous space can be challenging, and traversing the entire list may be required in some cases.

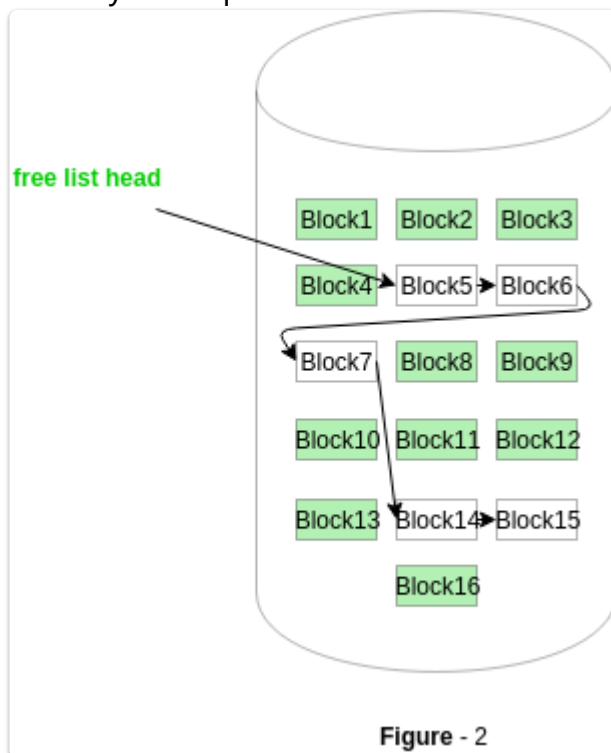
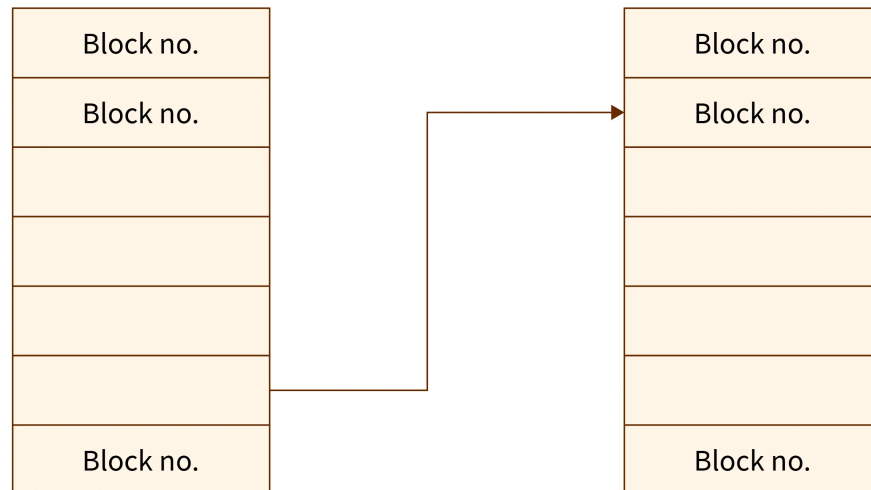


Figure - 2

3. Grouping:


- This method modifies the linked list approach to store the address of the next n-1 free blocks in the first free block.
- Additionally, it includes a pointer to the next block containing free-block-pointers.
- Grouping allows for easy access to addresses of a group of free disk blocks, improving efficiency.



SCALER
Topics

4. Counting:

- With counting, the space is frequently used and freed in contiguous allocation schemes like extents or clustering.
- It keeps track of the address of the first free block and the count of following free blocks.
- The free space list contains entries containing addresses and counts, making it efficient for managing contiguous free space.



SCALER
Topics

10. Discuss in detail the file allocation techniques?


File Allocation Techniques

File allocation techniques are methods used by file systems to manage and allocate disk space for storing files efficiently. There are several file allocation techniques, each with its own advantages and disadvantages. The main file allocation techniques include **contiguous allocation**, **linked allocation**, **indexed allocation**, and variations like **File Allocation Table (FAT)** and **Virtual File Allocation Table (VFAT)**. Let's discuss each of these techniques in detail:

1. Contiguous Allocation:

- In **contiguous allocation**, each file occupies a set of contiguous disk blocks. This means that the blocks allocated to a file are physically adjacent to each other on the disk.
- Contiguous allocation offers good performance since reading and writing files is straightforward.
- It requires only the starting location (block #) and the length (number of blocks) to locate a file.
- However, contiguous allocation suffers from several drawbacks:
 - Finding contiguous space for a file can be challenging, especially as the disk becomes fragmented.
 - Knowing the exact file size beforehand is necessary, which might not always be feasible.
 - Contiguous allocation can lead to external fragmentation, where free space becomes scattered across the disk, making it difficult to allocate large contiguous blocks.
 - Compaction may be required to defragment the disk, but this process often requires downtime.
- Despite its drawbacks, contiguous allocation remains one of the simplest and most efficient methods, especially for smaller file systems.

Contiguous Allocation




File Name	Start	Length	Allocated Blocks
file1.txt	0	4	0,1,2,3
sun.jpg	5	3	5,6,7
mov.mp4	9	3	9,10,11

Directory

SCALER
Topics

2. Linked Allocation:

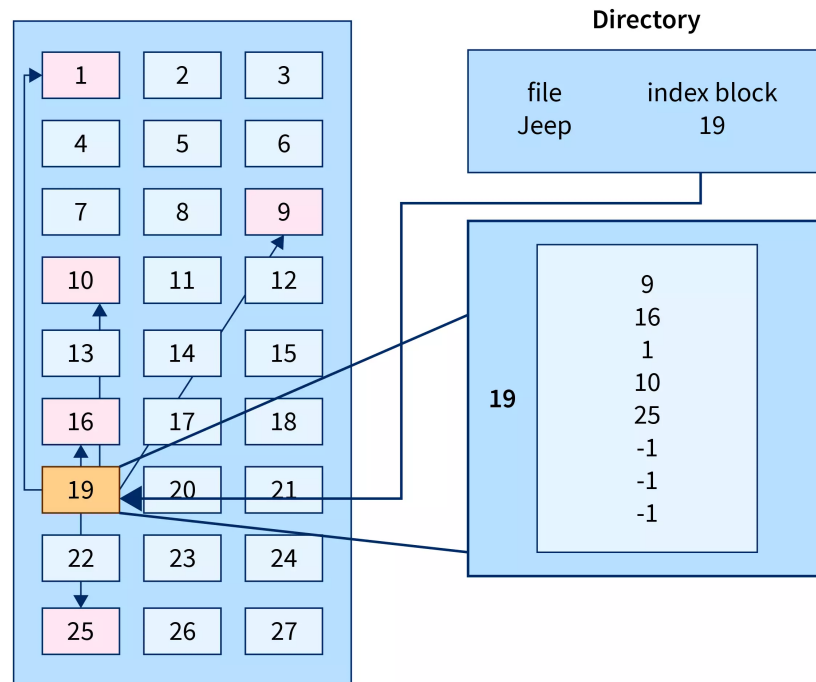
- **Linked allocation** represents each file as a linked list of disk blocks. Each block contains a pointer to the next block in the file.
- The file ends with a null pointer, indicating the end of the file.
- Linked allocation avoids external fragmentation since files can be allocated non-contiguously.
- However, linked allocation has its own set of disadvantages:
 - Locating a specific block within a file can be inefficient, as it may require multiple disk accesses and seeks.
 - The system needs to manage the pointers between blocks, which can add overhead.
 - Linked allocation can suffer from poor reliability since a single pointer failure can corrupt the entire file.
 - Free space management becomes crucial, as each block needs to be tracked individually.



3. Indexed Allocation:

- **Indexed allocation** maintains an index block for each file, containing pointers to the data blocks of the file.
- Unlike linked allocation, indexed allocation improves efficiency by clustering blocks into groups, reducing internal fragmentation.
- Each file has its own index block(s), which eliminates the need for traversing pointers to access blocks.
- Indexed allocation facilitates faster access to data blocks compared to linked allocation.
- However, indexed allocation may suffer from wasted space if the index blocks are not fully utilized.
- Implementations like the **File Allocation Table (FAT)** variation and **Virtual File Allocation Table (VFAT)** use indexed allocation techniques to manage file systems

efficiently.



SCALER
Topics

11. With neat diagram explain logical address and physical address?

Logical and Physical Addresses

Logical address and physical address are terms used in computer systems to describe different aspects of memory addressing:

1. Logical Address:

- A **logical address**, also known as a virtual address, is an address generated by the CPU during the execution of a program.
- It represents a location in the logical address space of a process.
- The logical address space is the set of all addresses that a process can reference or generate.
- The user program deals with logical addresses, which are typically in the range from 0 to some maximum value.
- Logical addresses are used by the CPU and the program during its execution. The program assumes that it has access to a contiguous block of memory starting

from address 0.

2. Physical Address:

- A **physical address** refers to the actual location in the physical memory (RAM) where data is stored.
- It represents a location in the physical address space of the system.
- Physical addresses are the addresses that the memory unit or the hardware recognizes and uses to access data.
- The physical address space is the set of all addresses corresponding to the physical memory installed in the system.
- The mapping from logical addresses to physical addresses is done by the memory management unit (MMU), a hardware component.
- The MMU uses techniques such as base and limit registers or page tables to translate logical addresses to physical addresses.
- The user program never directly interacts with physical addresses. It only deals with logical addresses, leaving the translation to the MMU.

